



**Updated for
Airflow 3**



The screenshot shows the Apache Airflow 3.0.0 web interface. The top navigation bar includes links for Home, Dags, Assets, Browse, Admin, Docs, and User. The main content area displays the following information:

- Stats:** Failed Dags (0), Running Dags (1), Active Dags (4).
- Dags:** First 10 favorite Dags: ai_support_ticket_system, dag_versioning_example, s3_to_snowflake_pipeline.
- Health:** MetaDatabase, Scheduler, Triggerer, Dag Processor. Pool Slots: 128.
- History:** Last 24 Hours: 2025-09-29 11:24:45 - 2025-09-30 11:24:45. Dag Runs: Queued (0, 0.00%), Running (1, 2.86%), Success (34, 97.14%), Failed (0, 0.00%).
- Asset Events:** Newest First. Log entries:
 - 2025-09-30 11:24:02: model_trained, Source: dag_versioning_example
 - 2025-09-30 11:24:02: model_trained, Source: dag_versioning_example
 - 2025-09-30 11:24:01: model_trained, Source: dag_versioning_example
 - 2025-09-30 11:24:01: model_trained, Source: dag_versioning_example

ETL/ELT patterns with Apache Airflow® 3

9 Practical DAG Code Examples

ASTRONOMER

In this new version of this eBook you'll find ETL and ELT pattern DAGs updated for Airflow 3, as well as an ETL and an ELT sequence using the new `@asset` decorator!

PART ONE

DAG Code Examples 4

1. ETL pattern DAG using XCom	5
2. ETL pattern DAG using explicit external storage	11
3. ELT pattern DAG using XCom	19
4. ELT pattern DAG using explicit external storage	25
5. ETLT DAG using XCom	32
6. Advanced ETL DAG with data quality checks	40
7. ELT DAG using Dynamic Task Mapping	50

PART TWO

Asset Code Examples 57

8. Asset sequence ETL with XCom	60
9. Asset sequence ETL using explicit external storage	67

Introduction

When writing data pipelines, regardless of the end use case, whether you provide data for analytics, complex dashboards, or are fine-tuning large language models, your data must first be **extracted (E)** from its source, **transformed (T)** into a target schema, and **loaded (L)** into your central data storage, whether that's a data warehouse or a data lake.

ETL and ELT pipelines form the core foundation of any data architecture, making the ability to design, build, and understand these pipelines essential for every data engineer.

Originally developed in 2014 specifically for ETL and ELT pipelines, [Apache Airflow®](#) has since evolved into the open-source standard for full-stack orchestration of any workflow—without losing sight of its roots in moving and transforming data.

According to the [2025 State of Airflow Report](#), over 85% of current Airflow users rely on it to run ETL and ELT pipelines, powering everything from some of the world's largest organizations to fast-growing startups.

This eBook contains full example DAG patterns for different ETL and ELT scenarios. The code to all these DAGs can also be found in [this GitHub repository](#), which is set up for you to be able to run them locally without needing to configure any connections or additional tools.

While the example DAGs use Postgres as their target system, and MinIO as external storage, they can be adapted to interact with any tool that is accessible through an API. For many tools, pre-built modules exist, see the [Astronomer Registry](#) for a comprehensive list.

This eBook is aimed at intermediate Airflow users. If you are new to Airflow we recommend first completing the [Apache Airflow® Crash Course: From 0 to running your pipeline in the Cloud](#) to get familiar with core Airflow concepts before using the code in this eBook.

If you are looking for more guidance on how to use Airflow for ETL and ELT, check out the [Best practices for ELT and ETL pipelines with Apache Airflow® eBook](#) for all you need to know to write ETL and ELT pipelines.

PART ONE

DAG Code Examples

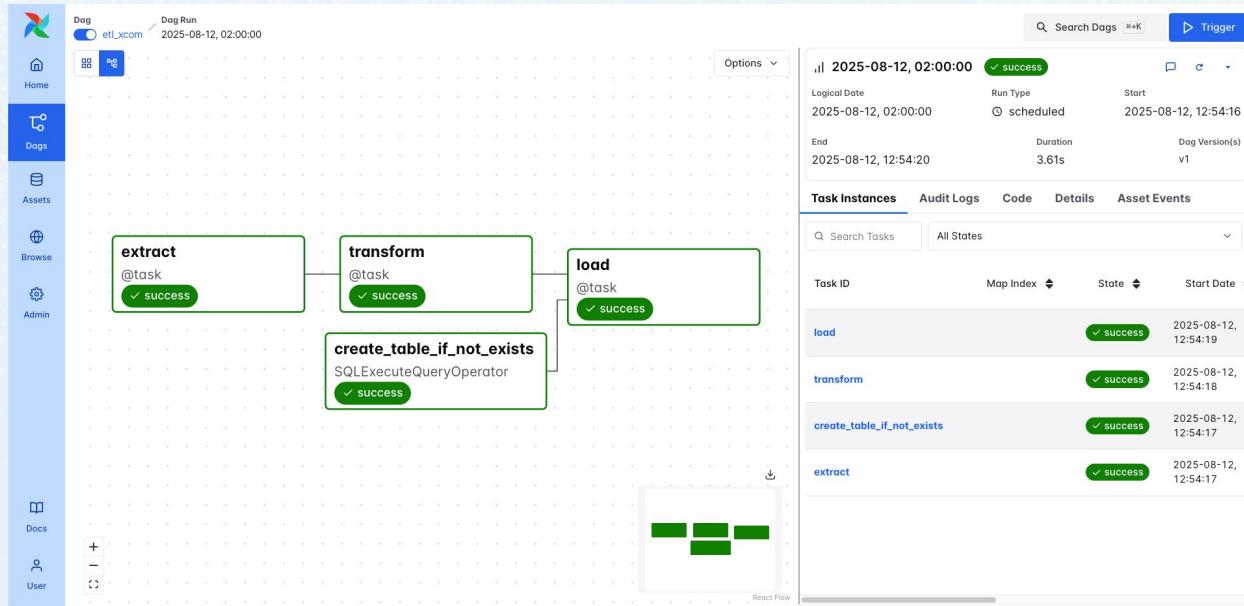
1

ETL pattern DAG using XCom

This DAG shows an ETL pattern to extract data from the Open-Meteo API, transform it and then load it into a Postgres database. The data is passed through XCom between the extraction and transformation task, as well as between the transformation and loading task, which means this setup would need a [custom XCom backend](#) in production.

In the companion GitHub repository you can find this DAG [here](#) and its supporting SQL scripts [here](#).

This is the graph of the DAG that will be created:



```
"""
## Simple ETL DAG loading data from the Open-Meteo API to a Postgres database
```

This DAG extracts weather data from the Open-Meteo API, transforms it, and loads it into a Postgres database in an ETL pattern.

It passes the data through XComs.

```
"""
```

```
import os
from datetime import datetime, timedelta

from airflow.sdk import dag, task, chain, Param
from airflow.providers.common.sql.operators.sql import SQLExecuteQueryOperator
from include.col_orders import WEATHER_COL_ORDER
```

```
# ----- #
# DAG-level variables #
# ----- #
```

```
DAG_ID = os.path.basename(__file__).replace(".py", "")
```

```
_POSTGRES_CONN_ID = os.getenv("POSTGRES_CONN_ID", "postgres_default")
_POSTGRES_DATABASE = os.getenv("POSTGRES_DATABASE", "postgres")
_POSTGRES_SCHEMA = os.getenv("POSTGRES_SCHEMA", "public")
_POSTGRES_TRANSFORMED_TABLE = os.getenv(
    "POSTGRES_WEATHER_TABLE_TRANSFORMED", f"model_weather_data_{DAG_ID}"
)
```

```
_SQL_DIR = f"{os.getenv('AIRFLOW_HOME')}/include/sql/pattern_dags/{DAG_ID}"
```

```
# ----- #
# DAG definition #
```

```

# -----
# @dag(
    dag_id=DAG_ID,
    start_date=datetime(2025, 8, 1), # date after which the DAG can be scheduled
    schedule="@daily", # see: https://www.astronomer.io/docs/learn/scheduling-in-airflow-for-options
    max_active_runs=1, # maximum number of active DAG runs
    max_consecutive_failed_dag_runs=5, # auto-pauses the DAG after 5 consecutive failed runs, experimental
    doc_md=__doc__, # add DAG Docs in the UI, see https://www.astronomer.io/docs/learn/custom-airflow-ui-docs-tutorial
    default_args={
        "owner": "Astro", # owner of this DAG in the Airflow UI
        "retries": 3, # tasks retry 3 times before they fail
        "retry_delay": timedelta(seconds=30), # tasks wait 30s in between retries
    },
    tags=["Patterns", "ETL", "Postgres", "XCom"], # add tags in the UI
    params={
        "coordinates": Param({"latitude": 46.9481, "longitude": 7.4474}, type="object")
    }, # Airflow params can add interactive options on manual runs. See: https://www.astronomer.io/docs/learn/airflow-params
    template_searchpath=[_SQL_DIR], # path to the SQL templates
)
def etl_xcom():

# -----
# Task Definitions #
# -----
# the @task decorator turns any Python function into an Airflow task
# any @task decorated function that is called inside the @dag decorated
# function is automatically added to the DAG.
# if one exists for your use case you can use traditional Airflow operators
# and mix them with @task decorators. Checkout registry.astronomer.io for available
operators
# see: https://www.astronomer.io/docs/learn/airflow-decorators for information about
@task
# see: https://www.astronomer.io/docs/learn/what-is-an-operator for information
about traditional operators

```

```
_create_table_if_not_exists = SQLExecuteQueryOperator(  
    task_id="create_table_if_not_exists",  
    conn_id=_POSTGRES_CONN_ID,  
    database=_POSTGRES_DATABASE,  
    sql="create_table_if_not_exists.sql",  
    params={"schema": _POSTGRES_SCHEMA, "table": _POSTGRES_TRANSFORMED_TABLE},  
)  
  
@task  
  
def extract(**context):  
    """  
  
    Extract data from the Open-Meteo API  
  
    Returns:  
  
        dict: The full API response  
  
    """  
  
    import requests  
  
    url = os.getenv("WEATHER_API_URL")  
  
    coordinates = context["params"]["coordinates"]  
    latitude = coordinates["latitude"]  
    longitude = coordinates["longitude"]  
  
    url = url.format(latitude=latitude, longitude=longitude)  
  
    response = requests.get(url)  
  
    return response.json()  
  
@task  
  
def transform(api_response: dict, **context) -> dict:  
    """  
  
    Transform the data  
  
    Args:  
  
        api_response (dict): The full API response  
  
    Returns:  
  
        dict: The transformed data
```

```
"""
time = api_response["hourly"]["time"]
dag_run_timestamp = context["ts"]

transformed_data = {
    "temperature_2m": api_response["hourly"]["temperature_2m"],
    "relative_humidity_2m": api_response["hourly"]["relative_humidity_2m"],
    "precipitation_probability": api_response["hourly"][
        "precipitation_probability"
    ],
    "timestamp": time,
    "date": [
        datetime.strptime(x, "%Y-%m-%dT%H:%M").date().strftime("%Y-%m-%d")
        for x in time
    ],
    "day": [datetime.strptime(x, "%Y-%m-%dT%H:%M").day for x in time],
    "month": [datetime.strptime(x, "%Y-%m-%dT%H:%M").month for x in time],
    "year": [datetime.strptime(x, "%Y-%m-%dT%H:%M").year for x in time],
    "last_updated": [dag_run_timestamp for i in range(len(time))],
    "latitude": [api_response["latitude"] for i in range(len(time))],
    "longitude": [api_response["longitude"] for i in range(len(time))],
}

return transformed_data

@task
def load(transformed_data: dict):
    """
    Load the data to the destination without using a temporary CSV file.

    Args:
        transformed_data (dict): The transformed data
    """

    import csv
    import io
    from airflow.providers.postgres.hooks.postgres import PostgresHook
```

```
hook = PostgresHook(postgres_conn_id=_POSTGRES_CONN_ID)

csv_buffer = io.StringIO()
writer = csv.writer(csv_buffer)
writer.writerow(WEATHER_COL_ORDER)
rows = zip(*[transformed_data[col] for col in WEATHER_COL_ORDER])
writer.writerows(rows)

csv_buffer.seek(0)

with open(f" {_SQL_DIR}/copy_insert.sql") as f:
    sql = f.read()
sql = sql.replace("{schema}", _POSTGRES_SCHEMA)
sql = sql.replace("{table}", _POSTGRES_TRANSFORMED_TABLE)

conn = hook.get_conn()
cursor = conn.cursor()
cursor.copy_expert(sql=sql, file=csv_buffer)
conn.commit()
cursor.close()
conn.close()

_extract = extract()
_transform = transform(api_response=_extract)
_load = load(transformed_data=_transform)
chain(_transform, _load)
chain(_create_table_if_not_exists, _load)

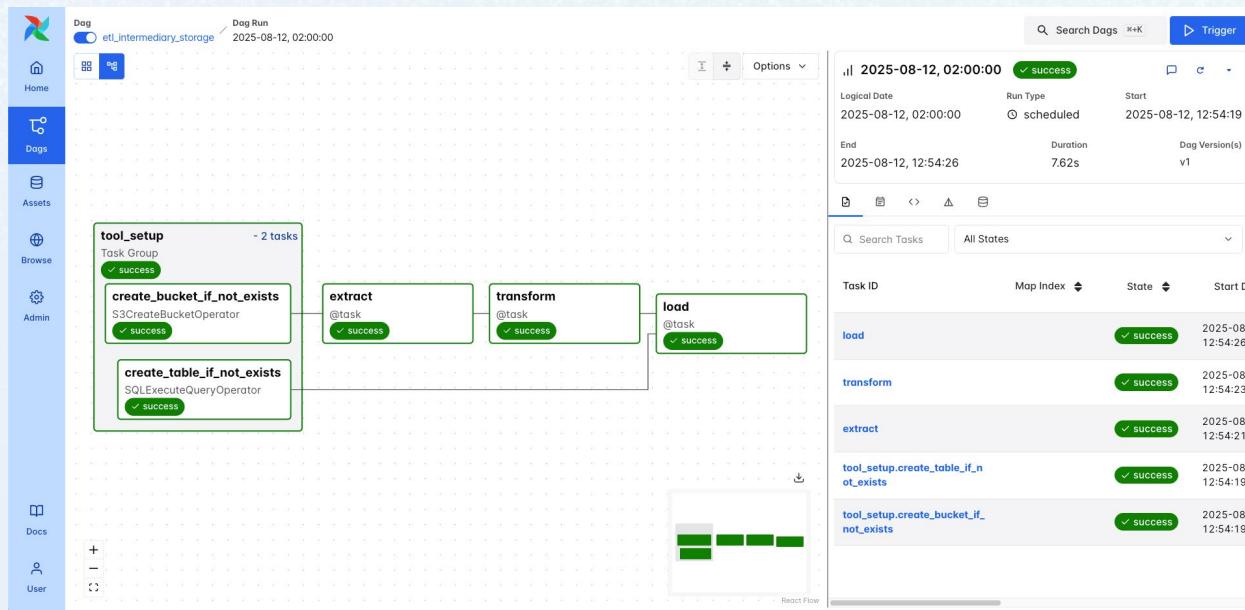
etl_xcom()
```

ETL pattern DAG using explicit external storage

This DAG accomplishes the same task, extracting data from the Open-Meteo API, transforming it and loading it to a Postgres database, but it does not pass any data through XCom, instead, MinIO is used as explicit external storage.

In the companion GitHub repository you can find this DAG [here](#) and its supporting SQL scripts [here](#).

This is the graph of the DAG that is created:



```
"""
## Simple ETL DAG loading data from the Open-Meteo API to a Postgres database
```

This DAG extracts weather data from the Open-Meteo API, transforms it, and loads it into a Postgres database in an ETL pattern.

It uses S3 as an intermediary storage.

```
import json
import os
from datetime import datetime, timedelta

import requests
from airflow.sdk import dag, task, task_group, chain, Param
from airflow.providers.amazon.aws.hooks.s3 import S3Hook
from airflow.providers.amazon.aws.operators.s3 import S3CreateBucketOperator
from airflow.providers.common.sql.operators.sql import SQLExecuteQueryOperator

# ----- #
# DAG-level variables #
# ----- #

DAG_ID = os.path.basename(__file__).replace(".py", "")

_AWS_CONN_ID = os.getenv("MINIO_CONN_ID", "minio_local")
_S3_BUCKET = os.getenv("S3_BUCKET", "open-meteo-etl")

_POSTGRES_CONN_ID = os.getenv("POSTGRES_CONN_ID", "postgres_default")
_POSTGRES_DATABASE = os.getenv("POSTGRES_DATABASE", "postgres")
_POSTGRES_SCHEMA = os.getenv("POSTGRES_SCHEMA", "public")
_POSTGRES_TRANSFORMED_TABLE = os.getenv(
    "POSTGRES_WEATHER_TABLE_TRANSFORMED", f"model_weather_data_{DAG_ID}"
)
)
```

```
_SQL_DIR = f"{os.getenv('AIRFLOW_HOME')}/include/sql/pattern_dags/{DAG_ID}"  
  
_EXTRACT_TASK_ID = "extract"  
_TRANSFORM_TASK_ID = "transform"  
  
# ----- #  
# DAG definition #  
# ----- #  
  
@dag(  
    dag_id=DAG_ID,  
    start_date=datetime(2025, 8, 1),  # date after which the DAG can be scheduled  
    schedule="@daily",  # see: https://www.astronomer.io/docs/learn/scheduling-in-airflow-for-options  
    max_active_runs=1,  # maximum number of active DAG runs  
    max_consecutive_failed_dag_runs=5,  # auto-pauses the DAG after 5 consecutive failed runs, experimental  
    doc_md=__doc__,  # add DAG Docs in the UI, see https://www.astronomer.io/docs/learn/custom-airflow-ui-docs-tutorial  
    default_args={  
        "owner": "Astro",  # owner of this DAG in the Airflow UI  
        "retries": 3,  # tasks retry 3 times before they fail  
        "retry_delay": timedelta(seconds=30),  # tasks wait 30s in between retries  
    },  
    tags=["Patterns", "ETL", "Postgres", "Intermediary Storage"],  # add tags in the UI  
    params={  
        "coordinates": Param({"latitude": 46.9481, "longitude": 7.4474}, type="object")  
    },  # Airflow params can add interactive options on manual runs. See: https://www.astronomer.io/docs/learn/airflow-params  
    template_searchpath=[_SQL_DIR],  # path to the SQL templates  
)  
  
def etl_intermediary_storage():  
  
    # ----- #  
    # Task Definitions #  
    # ----- #  
    # the @task decorator turns any Python function into an Airflow task
```

```
# any @task decorated function that is called inside the @dag decorated
# function is automatically added to the DAG.

# if one exists for your use case you can use traditional Airflow operators
# and mix them with @task decorators. Checkout registry.astronomer.io for available
operators

# see: https://www.astronomer.io/docs/learn/airflow-decorators for information about
@task

# see: https://www.astronomer.io/docs/learn/what-is-an-operator for information
about traditional operators

@task_group
def tool_setup():

    _create_table_if_not_exists = SQLExecuteQueryOperator(
        task_id="create_table_if_not_exists",
        conn_id=_POSTGRES_CONN_ID,
        database=_POSTGRES_DATABASE,
        sql="create_table_if_not_exists.sql",
        params={"schema": _POSTGRES_SCHEMA, "table": _POSTGRES_TRANSFORMED_TABLE},
    )

    _create_bucket_if_not_exists = S3CreateBucketOperator(
        task_id="create_bucket_if_not_exists",
        bucket_name=_S3_BUCKET,
        aws_conn_id=_AWS_CONN_ID,
    )

    return _create_table_if_not_exists, _create_bucket_if_not_exists

_tool_setup = tool_setup()

@task(task_id=_EXTRACT_TASK_ID)
def extract(**context):
    """
    Extract data from the Open-Meteo API
    Returns:
        dict: The full API response
    """

```

```
url = os.getenv("WEATHER_API_URL")

coordinates = context["params"]["coordinates"]
latitude = coordinates["latitude"]
longitude = coordinates["longitude"]
dag_run_timestamp = context["ts"]
dag_id = context["dag"].dag_id
task_id = context["task"].task_id

url = url.format(latitude=latitude, longitude=longitude)

response = requests.get(url).json()

response_bytes = json.dumps(response).encode("utf-8")

# Save the data to S3
hook = S3Hook(aws_conn_id=_AWS_CONN_ID)
hook.load_bytes(
    bytes_data=response_bytes,
    key=f"{dag_id}/{task_id}/{dag_run_timestamp}.json",
    bucket_name=_S3_BUCKET,
    replace=True,
)

@task(task_id=_TRANSFORM_TASK_ID)
def transform(**context):
    """
    Transform the data
    Args:
        api_response (dict): The full API response
    Returns:
        dict: The transformed data
    """
    dag_run_timestamp = context["ts"]
    dag_id = context["dag"].dag_id
```

```
upstream_task_id = _EXTRACT_TASK_ID
task_id = context["task"].task_id

# Load the data from S3
hook = S3Hook(aws_conn_id=_AWS_CONN_ID)
response = hook.read_key(
    key=f"{dag_id}/{upstream_task_id}/{dag_run_timestamp}.json",
    bucket_name=_S3_BUCKET,
)
api_response = json.loads(response)

time = api_response["hourly"]["time"]

transformed_data = {
    "temperature_2m": api_response["hourly"]["temperature_2m"],
    "relative_humidity_2m": api_response["hourly"]["relative_humidity_2m"],
    "precipitation_probability": api_response["hourly"][
        "precipitation_probability"
    ],
    "timestamp": time,
    "date": [
        datetime.strptime(x, "%Y-%m-%dT%H:%M").date().strftime("%Y-%m-%d")
        for x in time
    ],
    "day": [datetime.strptime(x, "%Y-%m-%dT%H:%M").day for x in time],
    "month": [datetime.strptime(x, "%Y-%m-%dT%H:%M").month for x in time],
    "year": [datetime.strptime(x, "%Y-%m-%dT%H:%M").year for x in time],
    "last_updated": [dag_run_timestamp for i in range(len(time))],
    "latitude": [api_response["latitude"] for i in range(len(time))],
    "longitude": [api_response["longitude"] for i in range(len(time))],
}

transformed_data_bytes = json.dumps(transformed_data).encode("utf-8")

# Save the data to S3
hook = S3Hook(aws_conn_id=_AWS_CONN_ID)
hook.load_bytes(
```

```
        bytes_data=transformed_data_bytes,
        key=f"{dag_id}/{task_id}/{dag_run_timestamp}.json",
        bucket_name=_S3_BUCKET,
        replace=True,
    )

@task
def load(**context):
    """
    Load the data to the destination without using a temporary CSV file.

    Args:
        transformed_data (dict): The transformed data
    """

    import csv
    import io

    from airflow.providers.postgres.hooks.postgres import PostgresHook

    dag_run_timestamp = context["ts"]
    dag_id = context["dag"].dag_id
    upstream_task_id = _TRANSFORM_TASK_ID

    # Load the data from S3
    hook = S3Hook(aws_conn_id=_AWS_CONN_ID)
    response = hook.read_key(
        key=f"{dag_id}/{upstream_task_id}/{dag_run_timestamp}.json",
        bucket_name=_S3_BUCKET,
    )
    api_response = json.loads(response)

    # Load the data to Postgres
    hook = PostgresHook(postgres_conn_id=_POSTGRES_CONN_ID)

    csv_buffer = io.StringIO()
    writer = csv.writer(csv_buffer)
    writer.writerow(api_response.keys())
    rows = zip(*api_response.values())
    writer.writerows(rows)
```

```
writer.writerows(rows)

csv_buffer.seek(0)

with open(f" {_SQL_DIR}/copy_insert.sql") as f:
    sql = f.read()
    sql = sql.replace("{schema}", _POSTGRES_SCHEMA)
    sql = sql.replace("{table}", _POSTGRES_TRANSFORMED_TABLE)

    conn = hook.get_conn()
    cursor = conn.cursor()
    cursor.copy_expert(sql=sql, file=csv_buffer)
    conn.commit()
    cursor.close()
    conn.close()

_extract = extract()
_transform = transform()
_load = load()
chain(_extract, _transform, _load)
chain(_tool_setup[0], _load)
chain(_tool_setup[1], _extract)

etl_intermediary_storage()
```

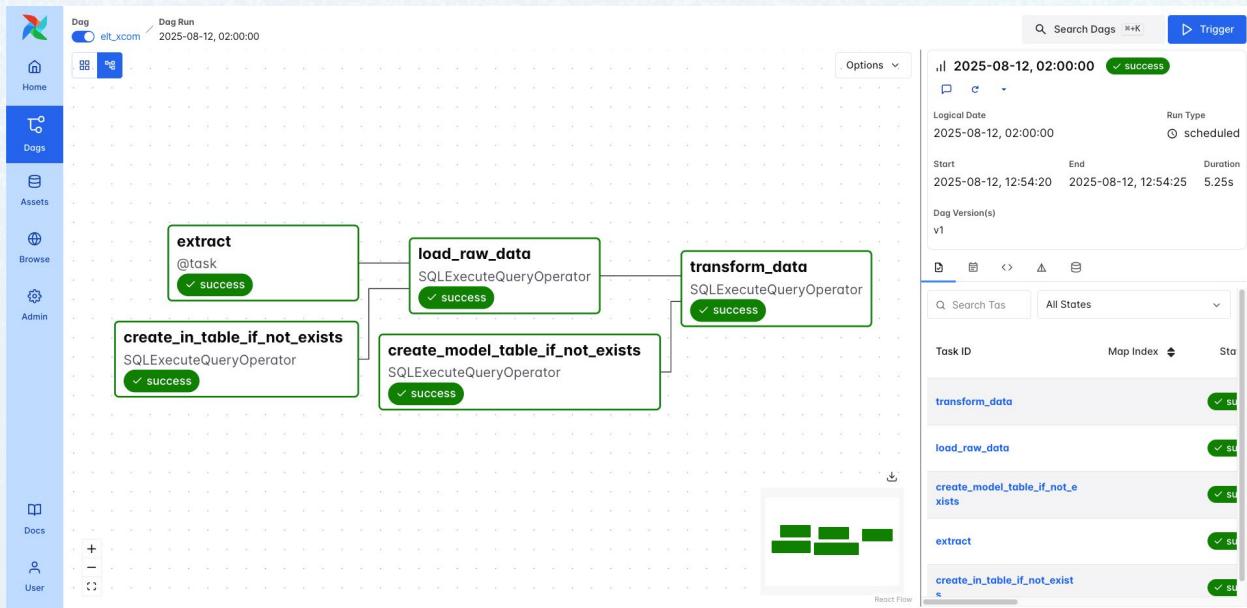
3

ELT pattern DAG using XCom

In this example DAG, the full data extracted from the Open-Meteo API is loaded into a Postgres database and all transformation happens inside the Postgres database using SQL. The data is passed through XCom between the extraction and loading task, which means this setup would need a [custom XCom backend](#) in production.

In the companion GitHub repository you can find this DAG [here](#) and its supporting SQL scripts [here](#).

This is the DAG that is being created:



```
"""

## Simple ELT DAG loading data from the Open-Meteo API to a Postgres database

This DAG extracts weather data from the Open-Meteo API,
loads it into a Postgres database and transforms it, using an ELT pattern.

It passes the data through XComs between extract and transform.

"""

import os

import json

from datetime import datetime, timedelta

from airflow.sdk import dag, task, chain, Param

from airflow.providers.common.sql.operators.sql import SQLExecuteQueryOperator

from include.col_orders import WEATHER_COL_ORDER


# ----- #
# DAG-level variables #
# ----- #


DAG_ID = os.path.basename(__file__).replace(".py", "")

_POSTGRES_CONN_ID = os.getenv("POSTGRES_CONN_ID", "postgres_default")
_POSTGRES_DATABASE = os.getenv("POSTGRES_DATABASE", "postgres")
_POSTGRES_SCHEMA = os.getenv("POSTGRES_SCHEMA", "public")
_POSTGRES_IN_TABLE = os.getenv("POSTGRES_WEATHER_TABLE_IN", f"in_weather_data_{DAG_ID}")
_POSTGRES_TRANSFORMED_TABLE = os.getenv(
    "POSTGRES_WEATHER_TABLE_TRANSFORMED", f"model_weather_data_{DAG_ID}"
```

```
)  
  
_SQL_DIR = f"{os.getenv('AIRFLOW_HOME')}/include/sql/pattern_dags/{DAG_ID}"  
  
# ----- #  
  
# DAG definition #  
  
# ----- #  
  
@dag(  
    dag_id=DAG_ID,  
    start_date=datetime(2025, 8, 1),  # date after which the DAG can be scheduled  
    schedule="@daily",  # see: https://www.astronomer.io/docs/learn/scheduling-in-airflow-for-options  
    max_active_runs=1,  # maximum number of active DAG runs  
    max_consecutive_failed_dag_runs=5,  # auto-pauses the DAG after 5 consecutive failed runs, experimental  
    doc_md=__doc__,  # add DAG Docs in the UI, see https://www.astronomer.io/docs/learn/custom-airflow-ui-docs-tutorial  
    default_args={  
        "owner": "Astro",  # owner of this DAG in the Airflow UI  
        "retries": 3,  # tasks retry 3 times before they fail  
        "retry_delay": timedelta(seconds=30),  # tasks wait 30s in between retries  
    },  
    tags=["Patterns", "ELT", "Postgres", "XCom"],  # add tags in the UI  
    params={  
        "coordinates": Param({"latitude": 46.9481, "longitude": 7.4474}, type="object")  
    },  # Airflow params can add interactive options on manual runs. See: https://www.astronomer.io/docs/learn/airflow-params
```

```
template_searchpath=[_SQL_DIR], # path to the SQL templates
)

def elt_xcom():

    # ----- #
    # Task Definitions #
    # ----- #

    # the @task decorator turns any Python function into an Airflow task
    # any @task decorated function that is called inside the @dag decorated
    # function is automatically added to the DAG.
    # if one exists for your use case you can use traditional Airflow operators
    # and mix them with @task decorators. Checkout registry.astronomer.io for available
    # operators
    # see: https://www.astronomer.io/docs/learn/airflow-decorators for information about
    # @task
    # see: https://www.astronomer.io/docs/learn/what-is-an-operator for information
    # about traditional operators

    _create_in_table_if_not_exists = SQLExecuteQueryOperator(
        task_id="create_in_table_if_not_exists",
        conn_id=_POSTGRES_CONN_ID,
        database=_POSTGRES_DATABASE,
        sql="create_in_table_if_not_exists.sql",
        params={"schema": _POSTGRES_SCHEMA, "table": _POSTGRES_IN_TABLE},
    )

    _create_model_table_if_not_exists = SQLExecuteQueryOperator(
        task_id="create_model_table_if_not_exists",
        conn_id=_POSTGRES_CONN_ID,
        database=_POSTGRES_DATABASE,
        sql="create_model_table_if_not_exists.sql",
    )
```

```
    params={"schema": _POSTGRES_SCHEMA, "table": _POSTGRES_TRANSFORMED_TABLE},  
)  
  
@task  
  
def extract(**context):  
  
    """  
  
    Extract data from the Open-Meteo API  
  
    Returns:  
  
        dict: The full API response  
  
    """  
  
    import requests  
  
  
    url = os.getenv("WEATHER_API_URL")  
  
  
    coordinates = context["params"]["coordinates"]  
    latitude = coordinates["latitude"]  
    longitude = coordinates["longitude"]  
  
  
    url = url.format(latitude=latitude, longitude=longitude)  
  
  
    response = requests.get(url)  
  
  
    return response.json()  
  
  
_extract = extract()  
  
  
_load = SQLExecuteQueryOperator(  
    task_id="load_raw_data",  
    conn_id=_POSTGRES_CONN_ID,  
    sql="""
```

```
INSERT INTO {{ params.schema }}.{{ params.table }} (raw_data)
VALUES ('{{ ti.xcom_pull(task_ids='extract') | tojson }}'::jsonb);
"""

params={

    "schema": _POSTGRES_SCHEMA,
    "table": _POSTGRES_IN_TABLE,
},
)

_transform = SQLExecuteQueryOperator(
    task_id="transform_data",
    conn_id=_POSTGRES_CONN_ID,
    sql="transform.sql",
    params={

        "schema": _POSTGRES_SCHEMA,
        "in_table": _POSTGRES_IN_TABLE,
        "out_table": _POSTGRES_TRANSFORMED_TABLE,
    },
)

chain([_create_in_table_if_not_exists, _extract], _load, _transform)
chain(_create_model_table_if_not_exists, _transform)

elt_xcom()
```

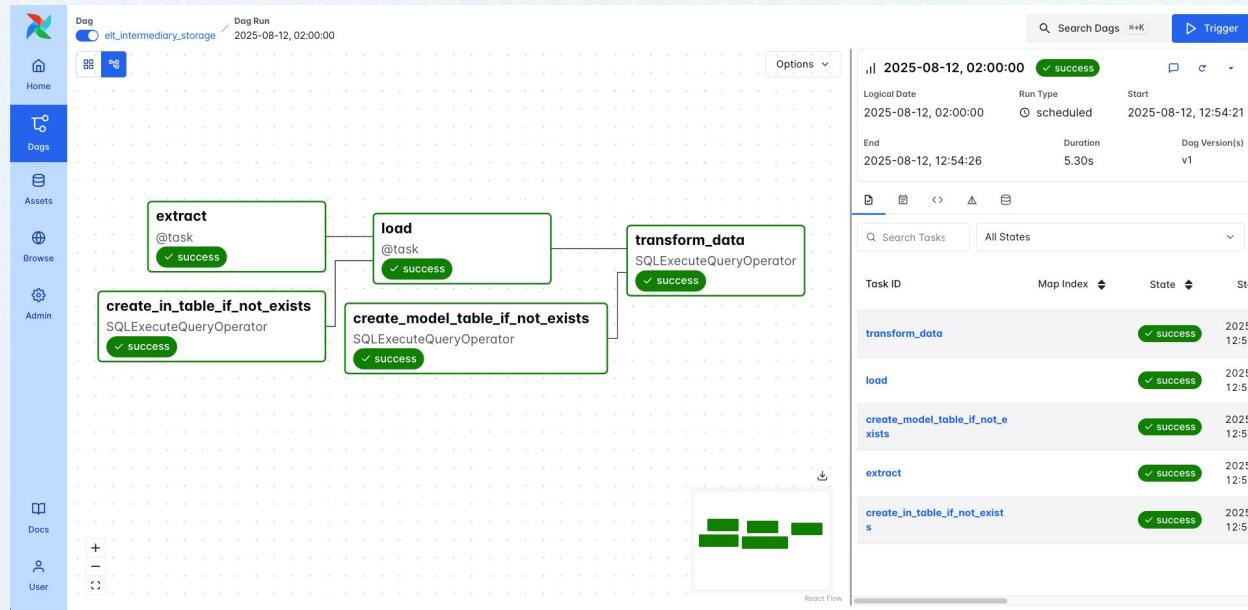
4

ELT pattern DAG using explicit external storage

This DAG accomplishes the same as the above DAG, extracting data from the Open-Meteo API, loading the raw data into a Postgres database and then running a transformation on it, but without passing any data through XCom.

In the companion GitHub repository you can find this DAG [here](#) and its supporting SQL scripts [here](#).

This is the DAG that will be created:



```
"""
```

```
## Simple ELT DAG loading data from the Open-Meteo API to a Postgres database
```

This DAG extracts weather data from the Open-Meteo API,

loads it into a Postgres database and transforms it, using an ELT pattern.

It passes the data through XCom between extract and transform.

```
"""
```

```
import os
import json
from datetime import datetime, timedelta
from airflow.providers.amazon.aws.hooks.s3 import S3Hook
from airflow.sdk import dag, task, chain, Param
from airflow.providers.common.sql.operators.sql import SQLExecuteQueryOperator
from airflow.providers.postgres.hooks.postgres import PostgresHook

# -----
# DAG-level variables #
# -----


DAG_ID = os.path.basename(__file__).replace(".py", "")

_AWS_CONN_ID = os.getenv("MINIO_CONN_ID", "minio_local")
_S3_BUCKET = os.getenv("S3_BUCKET", "open-meteo-etl")

_POSTGRES_CONN_ID = os.getenv("POSTGRES_CONN_ID", "postgres_default")
```

```
_POSTGRES_DATABASE = os.getenv("POSTGRES_DATABASE", "postgres")

_POSTGRES_SCHEMA = os.getenv("POSTGRES_SCHEMA", "public")

_POSTGRES_IN_TABLE = os.getenv("POSTGRES_WEATHER_TABLE_IN", f"in_weather_data_{DAG_ID}")

_POSTGRES_TRANSFORMED_TABLE = os.getenv(
    "POSTGRES_WEATHER_TABLE_TRANSFORMED", f"model_weather_data_{DAG_ID}"
)

_SQL_DIR = f"{os.getenv('AIRFLOW_HOME')}/include/sql/pattern_dags/{DAG_ID}"

_EXTRACT_TASK_ID = "extract"

# ----- #
# DAG definition #
# ----- #

@dag(
    dag_id=DAG_ID,
    start_date=datetime(2025, 8, 1), # date after which the DAG can be scheduled
    schedule="@daily", # see: https://www.astronomer.io/docs/learn/scheduling-in-airflow-for-options
    max_active_runs=1, # maximum number of active DAG runs
    max_consecutive_failed_dag_runs=5, # auto-pauses the DAG after 5 consecutive failed runs, experimental
    doc_md=__doc__, # add DAG Docs in the UI, see https://www.astronomer.io/docs/learn/custom-airflow-ui-docs-tutorial
    default_args={

        "owner": "Astro", # owner of this DAG in the Airflow UI

        "retries": 3, # tasks retry 3 times before they fail

        "retry_delay": timedelta(seconds=30), # tasks wait 30s in between retries
    }
)
```

```
},  
tags=[ "Patterns", "ELT", "Postgres", "XCom"], # add tags in the UI  
params={  
    "coordinates": Param({ "latitude": 46.9481, "longitude": 7.4474}, type="object")  
}, # Airflow params can add interactive options on manual runs. See: https://www.astronomer.io/docs/learn/airflow-params  
template_searchpath=[_SQL_DIR], # path to the SQL templates  
)  
  
def elt_intermediary_storage():  
  
    # ----- #  
  
    # Task Definitions #  
  
    # ----- #  
  
    # the @task decorator turns any Python function into an Airflow task  
  
    # any @task decorated function that is called inside the @dag decorated  
  
    # function is automatically added to the DAG.  
  
    # if one exists for your use case you can use traditional Airflow operators  
  
    # and mix them with @task decorators. Checkout registry.astronomer.io for available  
operators  
  
    # see: https://www.astronomer.io/docs/learn/airflow-decorators for information about  
@task  
  
    # see: https://www.astronomer.io/docs/learn/what-is-an-operator for information  
about traditional operators  
  
  
    _create_in_table_if_not_exists = SQLExecuteQueryOperator(  
        task_id="create_in_table_if_not_exists",  
        conn_id=_POSTGRES_CONN_ID,  
        database=_POSTGRES_DATABASE,  
        sql="create_in_table_if_not_exists.sql",  
        params={"schema": _POSTGRES_SCHEMA, "table": _POSTGRES_IN_TABLE},  
    )
```

```
_create_model_table_if_not_exists = SQLExecuteQueryOperator(  
    task_id="create_model_table_if_not_exists",  
    conn_id=_POSTGRES_CONN_ID,  
    database=_POSTGRES_DATABASE,  
    sql="create_model_table_if_not_exists.sql",  
    params={"schema": _POSTGRES_SCHEMA, "table": _POSTGRES_TRANSFORMED_TABLE},  
)  
  
@task(task_id=_EXTRACT_TASK_ID)  
  
def extract(**context):  
    """  
    Extract data from the Open-Meteo API  
  
    Returns:  
        dict: The full API response  
    """  
  
    import requests  
  
    url = os.getenv("WEATHER_API_URL")  
  
    coordinates = context["params"]["coordinates"]  
    latitude = coordinates["latitude"]  
    longitude = coordinates["longitude"]  
    dag_run_timestamp = context["ts"]  
    dag_id = context["dag"].dag_id  
    task_id = context["task"].task_id  
  
    url = url.format(latitude=latitude, longitude=longitude)  
  
    response = requests.get(url).json()
```

```
response_bytes = json.dumps(response).encode("utf-8")  
  
# Save the data to S3  
  
hook = S3Hook(aws_conn_id=_AWS_CONN_ID)  
  
hook.load_bytes(  
  
    bytes_data=response_bytes,  
  
    key=f"{dag_id}/{task_id}/{dag_run_timestamp}.json",  
  
    bucket_name=_S3_BUCKET,  
  
    replace=True,  
  
)  
  
  
_extract = extract()  
  
  
@task  
  
def load(**context):  
  
    """  
  
    Load the data from S3 to Postgres  
  
    """  
  
  
    dag_run_timestamp = context["ts"]  
  
    dag_id = context["dag"].dag_id  
  
    upstream_task_id = _EXTRACT_TASK_ID  
  
  
  
    print(f"{dag_id}/{upstream_task_id}/{dag_run_timestamp}.json")  
  
  
s3_hook = S3Hook(aws_conn_id=_AWS_CONN_ID)  
  
response = s3_hook.read_key(  
  
    key=f"{dag_id}/{upstream_task_id}/{dag_run_timestamp}.json",  
  
    bucket_name=_S3_BUCKET,
```

```
)  
api_response = json.loads(response)  
  
postgres_hook = PostgresHook(postgres_conn_id=_POSTGRES_CONN_ID)  
  
insert_sql = f"""  
INSERT INTO {_POSTGRES_SCHEMA}.{_POSTGRES_IN_TABLE} (raw_data)  
VALUES (%s::jsonb);  
"""  
  
postgres_hook.run(insert_sql, parameters=(json.dumps(api_response),))  
  
_transform = SQLExecuteQueryOperator(  
    task_id="transform_data",  
    conn_id=_POSTGRES_CONN_ID,  
    sql="transform.sql",  
    params={  
        "schema": _POSTGRES_SCHEMA,  
        "in_table": _POSTGRES_IN_TABLE,  
        "out_table": _POSTGRES_TRANSFORMED_TABLE,  
    },  
)  
  
chain([_create_in_table_if_not_exists, _extract], load(), _transform)  
chain(_create_model_table_if_not_exists, _transform)  
  
elt_intermediary_storage()
```

ETLT DAG using XCom

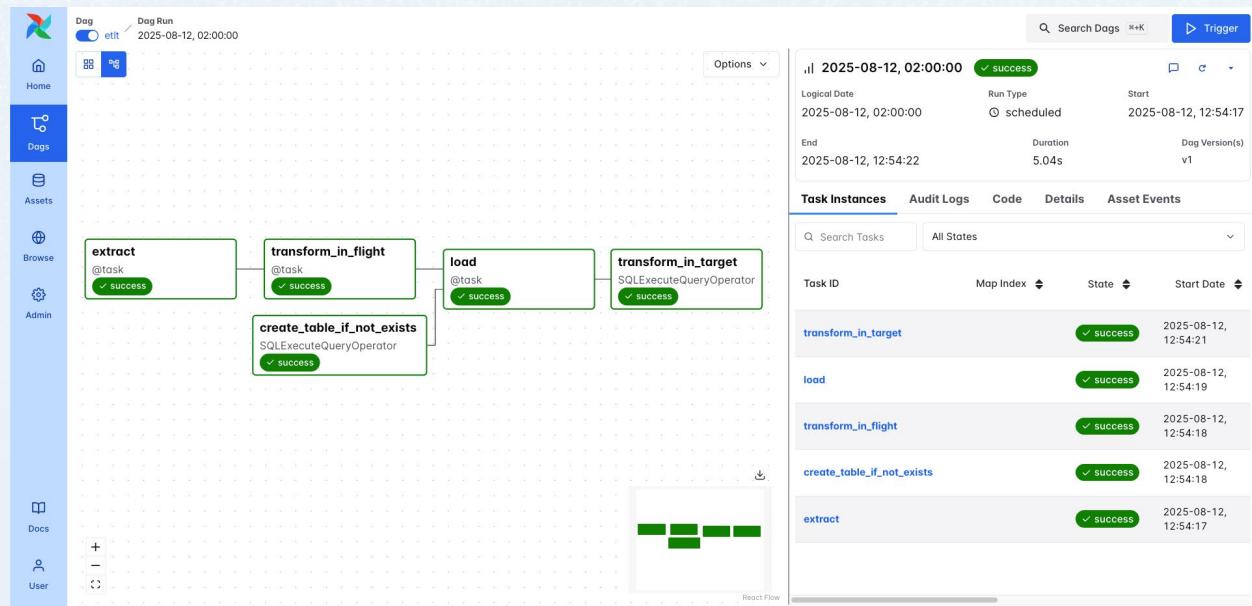
In real world scenarios ETL and ELT are often combined to form ETLT, a pattern where both types of transformations happen, once in flight before loading the data and once more after loading the data to a data warehouse.

This example shows the same ETL steps as the [ETL pattern DAG using XCom DAG](#), with an added additional transformation step that creates a secondary table with aggregated information.

Note that the data is passed through XCom between the extraction and transformation task, as well as between the transformation and loading task, which means this setup would need a [custom XCom backend](#) in production.

In the companion GitHub repository you can find this DAG [here](#) and its supporting SQL scripts [here](#).

This is the DAG that will be created:



```
"""
```

```
## Simple ETLT DAG loading data from the Open-Meteo API to a Postgres database
```

This DAG extracts weather data from the Open-Meteo API, transforms it, and loads it into a Postgres database, where a second transformation is applied.

Note that it uses XComs to pass data between tasks.

```
"""
```

```
import os

from datetime import datetime, timedelta

from airflow.sdk import dag, task, chain, Param

from airflow.providers.common.sql.operators.sql import SQLExecuteQueryOperator

from include.col_orders import WEATHER_COL_ORDER


# ----- #
# DAG-level variables #
# ----- #


DAG_ID = os.path.basename(__file__).replace(".py", "")

_POSTGRES_CONN_ID = os.getenv("POSTGRES_CONN_ID", "postgres_default")

_POSTGRES_DATABASE = os.getenv("POSTGRES_DATABASE", "postgres")

_POSTGRES_SCHEMA = os.getenv("POSTGRES_SCHEMA", "public")

_POSTGRES_TRANSFORMED_TABLE = os.getenv(
    "POSTGRES_WEATHER_TABLE_TRANSFORMED", f"model_weather_data_{DAG_ID}"
)

)
```

5

```
_POSTGRES_SECONDARY_TABLE = os.getenv("POSTGRES_WEATHER_TABLE_SECONDARY", f"secondary_weather_data_{DAG_ID}")

_SQL_DIR = f"{os.getenv('AIRFLOW_HOME')}/include/sql/pattern_dags/{DAG_ID}"

# ----- #
# DAG definition #
# ----- #

@dag(
    dag_id=DAG_ID,
    start_date=datetime(2025, 8, 1), # date after which the DAG can be scheduled
    schedule="@daily", # see: https://www.astronomer.io/docs/learn/scheduling-in-airflow-for-options
    max_active_runs=1, # maximum number of active DAG runs
    max_consecutive_failed_dag_runs=5, # auto-pauses the DAG after 5 consecutive failed runs, experimental
    doc_md=__doc__, # add DAG Docs in the UI, see https://www.astronomer.io/docs/learn/custom-airflow-ui-docs-tutorial
    default_args={
        "owner": "Astro", # owner of this DAG in the Airflow UI
        "retries": 3, # tasks retry 3 times before they fail
        "retry_delay": timedelta(seconds=30), # tasks wait 30s in between retries
    },
    tags=["Patterns", "ETL", "Postgres", "XCom"], # add tags in the UI
    params={
        "coordinates": Param({"latitude": 46.9481, "longitude": 7.4474}, type="object")
    }, # Airflow params can add interactive options on manual runs.
```

See: <https://www.astronomer.io/docs/learn/airflow-params>

```
template_searchpath=[_SQL_DIR], # path to the SQL templates

)

def etlt():

# ----- #

# Task Definitions #

# ----- #

# the @task decorator turns any Python function into an Airflow task

# any @task decorated function that is called inside the @dag decorated

# function is automatically added to the DAG.

# if one exists for your use case you can use traditional Airflow operators

# and mix them with @task decorators. Checkout registry.astronomer.io for available

operators

# see: https://www.astronomer.io/docs/learn/airflow-decorators for information about

@task

# see: https://www.astronomer.io/docs/learn/what-is-an-operator for information

about traditional operators

_create_table_if_not_exists = SQLExecuteQueryOperator(

    task_id="create_table_if_not_exists",

    conn_id=_POSTGRES_CONN_ID,

    database=_POSTGRES_DATABASE,

    sql="create_table_if_not_exists.sql",

    params={"schema": _POSTGRES_SCHEMA, "table": _POSTGRES_TRANSFORMED_TABLE},

)

@task

def extract(**context):

    """

    Extract data from the Open-Meteo API
```

```
    Returns:  
        dict: The full API response  
  
    """  
  
    import requests  
  
  
    url = os.getenv("WEATHER_API_URL")  
  
  
    coordinates = context["params"]["coordinates"]  
    latitude = coordinates["latitude"]  
    longitude = coordinates["longitude"]  
  
  
    url = url.format(latitude=latitude, longitude=longitude)  
  
  
    response = requests.get(url)  
  
  
    return response.json()  
  
  
@task  
  
def transform_in_flight(api_response: dict, **context) -> dict:  
  
    """  
  
    Transform the data  
  
    Args:  
        api_response (dict): The full API response  
  
    Returns:  
        dict: The transformed data  
  
    """  
  
  
    time = api_response["hourly"]["time"]  
    dag_run_timestamp = context["ts"]
```

```
transformed_data = {

    "temperature_2m": api_response["hourly"]["temperature_2m"],

    "relative_humidity_2m": api_response["hourly"]["relative_humidity_2m"],

    "precipitation_probability": api_response["hourly"][
        "precipitation_probability"
    ],

    "timestamp": time,

    "date": [
        datetime.strptime(x, "%Y-%m-%dT%H:%M").date().strftime("%Y-%m-%d")
            for x in time
    ],
    "day": [datetime.strptime(x, "%Y-%m-%dT%H:%M").day for x in time],
    "month": [datetime.strptime(x, "%Y-%m-%dT%H:%M").month for x in time],
    "year": [datetime.strptime(x, "%Y-%m-%dT%H:%M").year for x in time],
    "last_updated": [dag_run_timestamp for i in range(len(time))],
    "latitude": [api_response["latitude"] for i in range(len(time))],
    "longitude": [api_response["longitude"] for i in range(len(time))],
}

return transformed_data
```

```
@task
```

```
def load(transformed_data: dict):
```

```
    """
```

```
    Load the data to the destination without using a temporary CSV file.
```

```
Args:
```

```
    transformed_data (dict): The transformed data
```

```
    """
```

```
import csv
import io

from airflow.providers.postgres.hooks.postgres import PostgresHook

hook = PostgresHook(postgres_conn_id=_POSTGRES_CONN_ID)

csv_buffer = io.StringIO()

writer = csv.writer(csv_buffer)

writer.writerow(WEATHER_COL_ORDER)

rows = zip(*[transformed_data[col] for col in WEATHER_COL_ORDER])

writer.writerows(rows)

csv_buffer.seek(0)

with open(f" {_SQL_DIR}/copy_insert.sql") as f:

    sql = f.read()

    sql = sql.replace("{schema}", _POSTGRES_SCHEMA)

    sql = sql.replace("{table}", _POSTGRES_TRANSFORMED_TABLE)

conn = hook.get_conn()

cursor = conn.cursor()

cursor.copy_expert(sql=sql, file=csv_buffer)

conn.commit()

cursor.close()

conn.close()

_transform_in_target = SQLExecuteQueryOperator(
    task_id="transform_in_target",
    conn_id=_POSTGRES_CONN_ID,
    sql="transform.sql",
```

```
params={  
    "schema": _POSTGRES_SCHEMA,  
    "in_table": _POSTGRES_TRANSFORMED_TABLE,  
    "out_table": _POSTGRES_SECONDARY_TABLE,  
},  
)  
  
_extract = extract()  
_transform = transform_in_flight(api_response=_extract)  
_load = load(transformed_data=_transform)  
chain(_transform, _load)  
chain(_create_table_if_not_exists, _load, _transform_in_target)  
  
etlt()
```

Advanced ETL DAG with data quality checks

6

This advanced DAG has the same extraction and transforming steps as the ETL pattern DAG using [XCom](#) but it loads the data using a temporary table and includes several data quality checks.

In summary this DAG:

1. Extracts weather data from the Open-Meteo API
2. Transforms the data
3. Loads the data into a temporary table in Postgres
4. Performs stopping data quality checks on the temporary table
5. Swaps the temporary table with the permanent table (and creates a backup of the old permanent table)
6. Drops the temporary table
7. Performs stopping data quality checks on the new permanent table
8. Drops the backup table
9. Performs additional warning data quality checks on the new permanent table

This is the DAG that is created:



You can learn more about how to use the SQL check operators to write data quality checks [here](#).

In the companion GitHub repository you can find this DAG [here](#) and its supporting SQL scripts [here](#).

```
"""
## ETL with Data Quality Checks and Temporary Tables
```

This DAG will perform the following steps:

1. Extract weather data from the Open-Meteo API
2. Transform the data
3. Load the data into a temporary table in Postgres
4. Perform stopping data quality checks on the temporary table
5. Swap the temporary table with the permanent table (and creating a backup of the old permanent table)
6. Drop the temporary table
7. Perform stopping data quality checks on the new permanent table
8. Drop the backup table
9. Perform additional warning data quality checks on the new permanent table

```
"""
```

```
import os

from datetime import datetime, timedelta


from airflow.sdk import dag, task, task_group, chain, Param

from airflow.providers.common.sql.operators.sql import (
    SQLColumnCheckOperator,
    SQLExecuteQueryOperator,
    SQLTableCheckOperator,
)

from include.col_orders import WEATHER_COL_ORDER
```

```
DAG_ID = os.path.basename(__file__).replace(".py", "")  
  
_POSTGRES_CONN_ID = os.getenv("POSTGRES_CONN_ID", "postgres_default")  
  
_POSTGRES_DATABASE = os.getenv("POSTGRES_DATABASE", "postgres")  
  
_POSTGRES_SCHEMA = os.getenv("POSTGRES_SCHEMA", "public")  
  
_POSTGRES_TMP_TABLE = os.getenv(  
    "POSTGRES_WEATHER_TABLE_TMP", f"tmp_weather_data_{DAG_ID}"  
)  
  
_POSTGRES_BACKUP_TABLE = os.getenv(  
    "POSTGRES_WEATHER_TABLE_BACKUP", f"backup_weather_data_{DAG_ID}"  
)  
  
_POSTGRES_TRANSFORMED_TABLE = os.getenv(  
    "POSTGRES_WEATHER_TABLE_TRANSFORMED", f"model_weather_data_{DAG_ID}"  
)  
  
_SQL_DIR = f"{os.getenv('AIRFLOW_HOME')}/include/sql/pattern_dags/{DAG_ID}"  
  
@dag(  
    dag_id=DAG_ID,  
    start_date=datetime(2025, 8, 1), # date after which the DAG can be scheduled  
    schedule="@daily", # see: https://www.astronomer.io/docs/learn/scheduling-in-airflow-for-options  
    max_active_runs=1, # maximum number of active DAG runs  
    max_consecutive_failed_dag_runs=5, # auto-pauses the DAG after 5 consecutive failed runs, experimental  
    doc_md=__doc__, # add DAG Docs in the UI, see https://www.astronomer.io/docs/learn/custom-airflow-ui-docs-tutorial  
    default_args={  
        "owner": "Astro", # owner of this DAG in the Airflow UI  
        "retries": 3, # tasks retry 3 times before they fail
```

```
"retry_delay": timedelta(seconds=30),  # tasks wait 30s in between retries
    "postgres_conn_id": _POSTGRES_CONN_ID,
    "conn_id": _POSTGRES_CONN_ID,
},
tags=["Patterns", "ETL", "Postgres", "XCom"], # add tags in the UI
params={
    "coordinates": Param({"latitude": 46.9481, "longitude": 7.4474}, type="object")
}, # Airflow params can add interactive options on manual runs. See: https://www.astronomer.io/docs/learn/airflow-params
template_searchpath=[_SQL_DIR], # path to the SQL templates
)

def etl_dq_checks_tmp_tables():

    @task

    def extract(**context):
        """
        Extract data from the Open-Meteo API
        Returns:
            dict: The full API response
        """
        import requests

        url = os.getenv("WEATHER_API_URL")

        coordinates = context["params"]["coordinates"]
        latitude = coordinates["latitude"]
        longitude = coordinates["longitude"]

        url = url.format(latitude=latitude, longitude=longitude)
```

```
response = requests.get(url)

return response.json()

@task

def transform_in_flight(api_response: dict, **context) -> dict:

    """
    Transform the data

    Args:
        api_response (dict): The full API response

    Returns:
        dict: The transformed data

    """

time = api_response["hourly"]["time"]
dag_run_timestamp = context["ts"]

transformed_data = {
    "temperature_2m": api_response["hourly"]["temperature_2m"],
    "relative_humidity_2m": api_response["hourly"]["relative_humidity_2m"],
    "precipitation_probability": api_response["hourly"][
        "precipitation_probability"
    ],
    "timestamp": time,
    "date": [
        datetime.strptime(x, "%Y-%m-%dT%H:%M").date().strftime("%Y-%m-%d")
        for x in time
    ],
    "day": [datetime.strptime(x, "%Y-%m-%dT%H:%M").day for x in time],
    "month": [datetime.strptime(x, "%Y-%m-%dT%H:%M").month for x in time],
    "year": [datetime.strptime(x, "%Y-%m-%dT%H:%M").year for x in time],
}
```

```
"last_updated": [dag_run_timestamp for i in range(len(time))],  
"latitude": [api_response["latitude"] for i in range(len(time))],  
"longitude": [api_response["longitude"] for i in range(len(time))],  
}  
  
return transformed_data  
  
_extract = extract()  
_transform = transform_in_flight(_extract)  
  
@task_group  
def load():  
    _create_tmp_table_if_not_exists = SQLExecuteQueryOperator(  
        task_id="create_tmp_table_if_not_exists",  
        database=_POSTGRES_DATABASE,  
        sql="create_tmp_table_if_not_exists.sql",  
        params={"schema": _POSTGRES_SCHEMA, "table": _POSTGRES_TMP_TABLE},  
    )  
  
    @task  
    def load_to_tmp(transformed_data: dict):  
        """  
        Load the data to the destination without using a temporary CSV file.  
        Args:  
            transformed_data (dict): The transformed data  
        """  
        import csv  
        import io  
  
        from airflow.providers.postgres.hooks.postgres import PostgresHook
```

```
hook = PostgresHook(postgres_conn_id=_POSTGRES_CONN_ID)

csv_buffer = io.StringIO()

writer = csv.writer(csv_buffer)

writer.writerow(WEATHER_COL_ORDER)

rows = zip(*[transformed_data[col] for col in WEATHER_COL_ORDER])

writer.writerows(rows)

csv_buffer.seek(0)

with open(f"{_SQL_DIR}/copy_insert.sql") as f:

    sql = f.read()

    sql = sql.replace("{schema}", _POSTGRES_SCHEMA)

    sql = sql.replace("{table}", _POSTGRES_TMP_TABLE)

conn = hook.get_conn()

cursor = conn.cursor()

cursor.copy_expert(sql=sql, file=csv_buffer)

conn.commit()

cursor.close()

conn.close()

_load_to_tmp = load_to_tmp(_transform)

@task_group

def test_tmp_table():

    SQLTableCheckOperator(

        task_id="test_table",
```

```
retry_on_failure="True",

table=f"({_POSTGRES_SCHEMA}).{_POSTGRES_TMP_TABLE}",

checks={

    "row_count_check": {"check_statement": "COUNT(*) > 10"},

},

)

swap = SQLExecuteQueryOperator(

task_id="swap",

sql="swap.sql",

params={

    "schema": _POSTGRES_SCHEMA,

    "tmp_table": _POSTGRES_TMP_TABLE,

    "perm_table": _POSTGRES_TRANSFORMED_TABLE,

},

)

drop_tmp = SQLExecuteQueryOperator(

task_id="drop_tmp",

sql=f"DROP TABLE IF EXISTS {_POSTGRES_SCHEMA}.{_POSTGRES_TMP_TABLE};",

)

@task_group

def test_perm_table():

SQLColumnCheckOperator(

    task_id="test_cols",

    retry_on_failure="True",

    table=f"({_POSTGRES_SCHEMA}).{_POSTGRES_TRANSFORMED_TABLE}",

    column_mapping={

        "month": {"min": {"geq_to": 1}, "max": {"leq_to": 12}},

        "year": {"min": {"geq_to": 2010}, "max": {"leq_to": 2020}}
    }
)
```

```
        "day": {"min": {"geq_to": 1}, "max": {"leq_to": 31}},  
    },  
  
    accept_none="True",  
)  
  
  
SQLTableCheckOperator(  
  
    task_id="test_table",  
  
    retry_on_failure="True",  
  
    table=f"({_POSTGRES_SCHEMA}).{_POSTGRES_TRANSFORMED_TABLE}",  
  
    checks={  
  
        "row_count_check": {"check_statement": "COUNT(*) > 10"},  
  
    },  
  
)  
  
  
drop_backup = SQLExecuteQueryOperator(  
  
    task_id="drop_backup",  
  
    sql=f"DROP TABLE IF EXISTS {_POSTGRES_SCHEMA}.{_POSTGRES_BACKUP_TABLE};",  
  
)  
  
  
chain(  
  
    _create_tmp_table_if_not_exists,  
  
    _load_to_tmp,  
  
    test_tmp_table(),  
  
    swap,  
  
    [drop_tmp, test_perm_table()],  
  
    drop_backup,  
  
)  
  
  
@task_group  
  
def validate():
```

```
test_cols = SQLColumnCheckOperator(  
    task_id="test_cols",  
    retry_on_failure=True,  
    table=f"({_POSTGRES_SCHEMA}).{_POSTGRES_TRANSFORMED_TABLE}",  
    column_mapping={  
        "temperature_2m": {"min": {"geq_to": -10}, "max": {"leq_to": 50}},  
    },  
    accept_none=True,  
)  
  
@task(trigger_rule="all_done")  
def sql_check_done():  
    return "Additional data quality checks are done!"  
  
test_cols >> sql_check_done()  
  
swap >> validate()  
  
load()  
  
etl_dq_checks_tmp_tables()
```

7

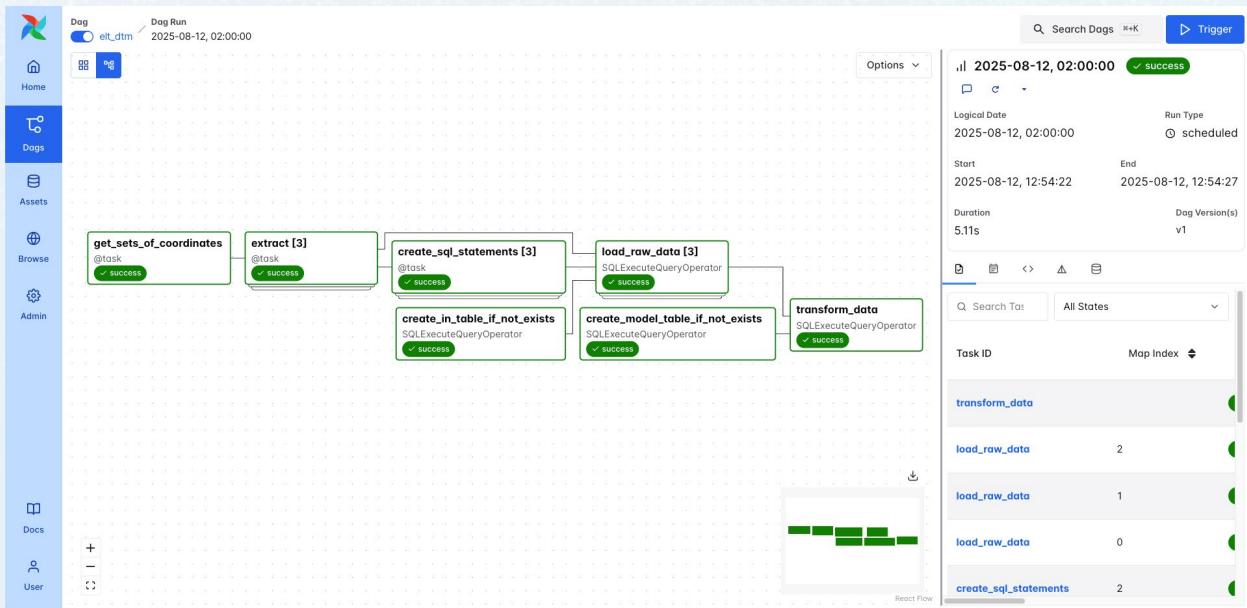
ELT DAG using Dynamic Task Mapping

The DAG below is similar to the [ELT pattern DAG using XCom](#), but instead of just fetching data from one API call including one set of coordinates, it uses [dynamic task mapping](#) to run 1 extraction and 1 loading tasks in parallel for each set of coordinates specified at runtime via the DAG configuration params.

The data is passed through XCom between the extraction and loading task, which means this setup would need a [custom XCom backend](#) in production.

In the companion GitHub repository you can find this DAG [here](#) and its supporting SQL scripts [here](#).

This is the DAG that will be created:



```
"""
```

```
## ELT DAG extracting data from 3 API calls in parallel using Dynamic Tasks
```

This DAG extracts weather data from the Open-Meteo API, loads it into a postgres database and transforms it, using an ELT pattern.

It extracts data from 3 different locations in parallel using dynamic task mapping.

```
"""
```

```
import json
import os
from datetime import datetime, timedelta

from airflow.sdk import dag, task, chain, Param
from airflow.providers.common.sql.operators.sql import SQLExecuteQueryOperator

# ----- #
# DAG-level variables #
# ----- #

DAG_ID = os.path.basename(__file__).replace(".py", "")

_POSTGRES_CONN_ID = os.getenv("POSTGRES_CONN_ID", "postgres_default")
_POSTGRES_DATABASE = os.getenv("POSTGRES_DATABASE", "postgres")
_POSTGRES_SCHEMA = os.getenv("POSTGRES_SCHEMA", "public")
_POSTGRES_IN_TABLE = os.getenv("POSTGRES_WEATHER_TABLE_IN", f"in_weather_data_{DAG_ID}")
_POSTGRES_TRANSFORMED_TABLE = os.getenv(
    "POSTGRES_WEATHER_TABLE_TRANSFORMED", f"model_weather_data_{DAG_ID}"
)
)
```

```
_SQL_DIR = f"{os.getenv('AIRFLOW_HOME')}/include/sql/pattern_dags/{DAG_ID}"
```

```
# ----- #
# DAG definition #
# ----- #

@dag(
    dag_id=DAG_ID,
    start_date=datetime(2025, 8, 1), # date after which the DAG can be scheduled
    schedule="@daily", # see: https://www.astronomer.io/docs/learn/scheduling-in-airflow-for-options
    max_active_runs=1, # maximum number of active DAG runs
    max_consecutive_failed_dag_runs=5, # auto-pauses the DAG after 5 consecutive failed runs, experimental
    doc_md=__doc__, # add DAG Docs in the UI, see https://www.astronomer.io/docs/learn/custom-airflow-ui-docs-tutorial
    default_args={
        "owner": "Astro", # owner of this DAG in the Airflow UI
        "retries": 3, # tasks retry 3 times before they fail
        "retry_delay": timedelta(seconds=30), # tasks wait 30s in between retries
    },
    tags=["Patterns", "ELT", "Postgres", "XCom"], # add tags in the UI
    params={
        "set_of_coordinates": Param(
            {
                "Location 1": {"latitude": 46.9480, "longitude": 7.4474},
                "Location 2": {"latitude": 38.4272, "longitude": 14.9524},
                "Location 3": {"latitude": 46.9480, "longitude": -2.9916},
            },
        ),
    },
)
```

```
        type="object",
    )

}, # Airflow params can add interactive options on manual runs. See: https://www.astronomer.io/docs/learn/airflow-params

template_searchpath=[_SQL_DIR], # path to the SQL templates

)

def elt_dtm():

# ----- #

# Task Definitions #

# ----- #

# the @task decorator turns any Python function into an Airflow task

# any @task decorated function that is called inside the @dag decorated

# function is automatically added to the DAG.

# if one exists for your use case you can use traditional Airflow operators

# and mix them with @task decorators. Checkout registry.astronomer.io for available

operators

# see: https://www.astronomer.io/docs/learn/airflow-decorators for information about

@task

# see: https://www.astronomer.io/docs/learn/what-is-an-operator for information

about traditional operators

_create_in_table_if_not_exists = SQLExecuteQueryOperator(

    task_id="create_in_table_if_not_exists",

    conn_id=_POSTGRES_CONN_ID,

    database=_POSTGRES_DATABASE,

    sql="create_in_table_if_not_exists.sql",

    params={"schema": _POSTGRES_SCHEMA, "table": _POSTGRES_IN_TABLE},

)

_create_model_table_if_not_exists = SQLExecuteQueryOperator(

    task_id="create_model_table_if_not_exists",
```

```
conn_id=_POSTGRES_CONN_ID,  
database=_POSTGRES_DATABASE,  
sql="create_model_table_if_not_exists.sql",  
params={"schema": _POSTGRES_SCHEMA, "table": _POSTGRES_TRANSFORMED_TABLE},  
)  
  
@task  
  
def get_sets_of_coordinates(**context):  
    """  
    Get the set of coordinates from the params  
    """  
  
    coordinate_dict = context["params"]["set_of_coordinates"]  
    list_of_coordinates = []  
  
    for k, v in coordinate_dict.items():  
        list_of_coordinates.append(v)  
  
    return list_of_coordinates  
  
_get_sets_of_coordinates = get_sets_of_coordinates()  
  
@task(map_index_template="{{ my_custom_map_index }}")  
def extract(coordinates: dict):  
    """  
    Extract data from the Open-Meteo API  
    Returns:  
        dict: The full API response  
    """  
  
    import requests  
  
    url = os.getenv("WEATHER_API_URL")
```

```
latitude = coordinates["latitude"]
longitude = coordinates["longitude"]

url = url.format(latitude=latitude, longitude=longitude)

response = requests.get(url)

# optional custom map index (2.9+)
from airflow.sdk import get_current_context

context = get_current_context()
context["my_custom_map_index"] = (
    "Coordinates: " + str(latitude) + ", " + str(longitude)
)

return response.json()

_extract = extract.expand(coordinates=_get_sets_of_coordinates)

@task
def create_sql_statements(data, schema, table):
    """
    Create SQL statements to load the data
    """
    data_json = json.dumps(data)

    return f"""
        INSERT INTO {schema}.{table} (raw_data)
        VALUES ('{data_json}'::jsonb);
    """
```

```
_create_sql_statements = create_sql_statements.partial(  
    schema=_POSTGRES_SCHEMA, table=_POSTGRES_IN_TABLE  
) .expand(data=_extract)  
  
_load = SQLExecuteQueryOperator.partial(  
    task_id="load_raw_data",  
    conn_id=_POSTGRES_CONN_ID,  
) .expand(sql=_create_sql_statements)  
  
_transform = SQLExecuteQueryOperator(  
    task_id="transform_data",  
    conn_id=_POSTGRES_CONN_ID,  
    sql="transform.sql",  
    params={  
        "schema": _POSTGRES_SCHEMA,  
        "in_table": _POSTGRES_IN_TABLE,  
        "out_table": _POSTGRES_TRANSFORMED_TABLE,  
    },  
)  
  
chain([_create_in_table_if_not_exists, _extract], _load, _transform)  
chain(_create_model_table_if_not_exists, _transform)  
  
elt_dtm()
```

PART TWO

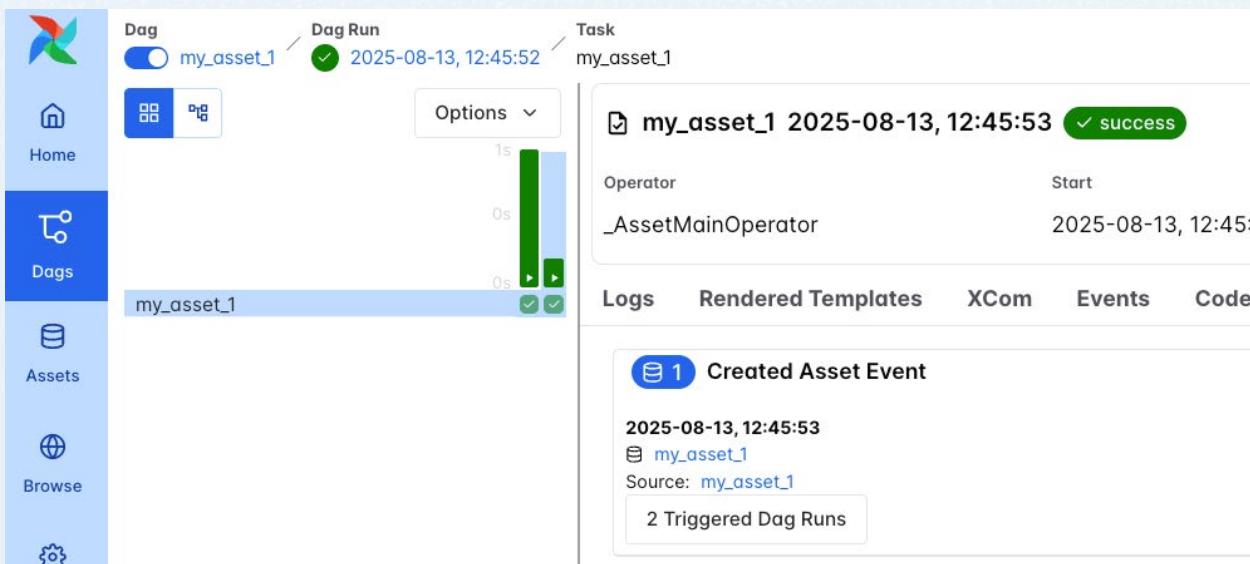
Asset Code Examples

Airflow 3 introduced an additional second way to write your pipeline: the asset-centric approach. You can use the `@asset` decorator as a shorthand to create one DAG containing one task that updates one Airflow Asset object.

```
from airflow.sdk import asset

@asset(schedule=None, tags=["basic_asset_example"])
def my_asset_1():
    print("Add any Python Code!")
```

The code above creates one DAG with the DAG ID `my_asset_1` containing one task with the task ID `my_asset_1` that produces updates to the Airflow Asset `my_asset_1`.



Assets can run on a schedule like regular DAGs, including data-aware schedules, i.e. an `@asset` can be scheduled based on updates to Assets. See the [Asset and data-aware scheduling guide](#) for more information.

It is also possible to push data to XCom from assets and retrieve them using a cross-DAG XCom pull, as well as adjust small amounts of `Metadata` to the extra dictionary of the Asset event.

The following code snippet shows several assets that depend on each other, push to and pull from XCom and use the extra dictionary of the Asset event. You can find the example [here on GitHub](#).

```
from airflow.sdk import asset, Metadata, Asset

@asset(schedule=None, tags=["basic_asset_example"])
def my_asset_1():
    print("Add any Python Code!")

@asset(schedule=[my_asset_1], tags=["basic_asset_example"])
def my_asset_2():
    # you can return data from an @asset decorated function
    # to push it to XCom
    return {"a": 1, "b": 2}

@asset(schedule=[my_asset_1, my_asset_2], tags=["basic_asset_example"])
def my_asset_3(context):
    # you can pull the data from XCom via the context
    # note that this is a cross-dag xcom pull
    data = context["ti"].xcom_pull(
        dag_id="my_asset_2",
        task_ids=["my_asset_2"],
        key="return_value",
        include_prior_dates=True,
    )[0]

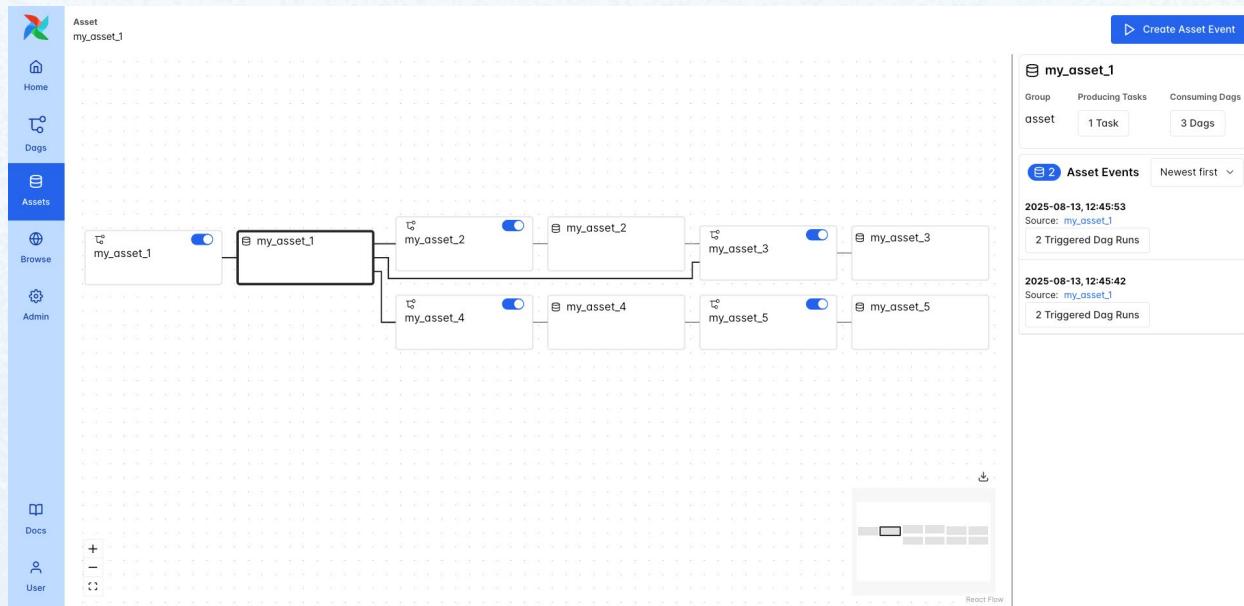
    print(data)

@asset(schedule=[my_asset_1], tags=["basic_asset_example"])
def my_asset_4(context):
```

```
# you can also attach Metadata to the Asset Event
yield Metadata(Asset("my_asset_4"), {"a": 1, "b": 2})
```

```
@asset(schedule=[my_asset_4], tags=["basic_asset_example"])
def my_asset_5(context):
    # and retrieve the Metadata from the Asset Event
    metadata = context["triggering_asset_events"][Asset("my_asset_4")][0].extra
    print(metadata)
```

In the Airflow 3 UI you can view asset dependencies by selecting a graph from the list shown under the **Assets** button. The following asset graph is created by the example code above.



8

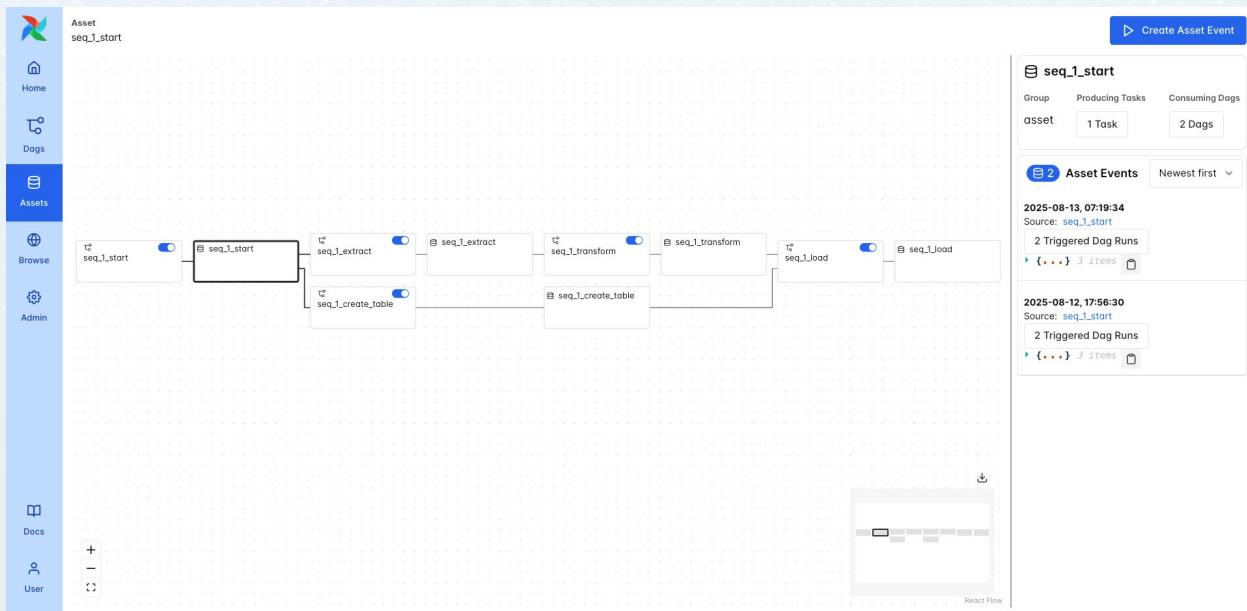
Asset sequence ETL with XCom

The sequence of assets below mirrors the [ETL pattern DAG using XCom](#) using the `@asset` decorator instead.

The sequence shows an ETL pattern to extract data from the Open-Meteo API, transform it and then load it into a Postgres database. The data is passed through XCom between the extraction and transformation task, as well as between the transformation and loading task, which means this setup would need a [custom XCom backend](#) in production.

In the companion GitHub repository you can find this DAG [here](#) and its supporting SQL scripts [here](#).

This is the asset graph that is created:



```
"""
## Simple asset sequence creating 4 DAGs to perform an ETL pattern

These 4 asset decorated functions created 4 DAGs each containing one task
following a similar pattern as the regular etl_xcom DAG.

Note that passing the data between assets via XCom necessitates a Cross-DAG xcom pull
and other information like DAG params and timestamps are passed via the metadata
asset events.

"""

import os
from datetime import datetime

from airflow.sdk import Param, asset, Metadata, Asset
from airflow.providers.postgres.hooks.postgres import PostgresHook
from include.col_orders import WEATHER_COL_ORDER

# ----- #
# DAG-level variables #
# ----- #

SEQUENCE_NAME = os.path.basename(__file__).replace(".py", "")

_POSTGRES_CONN_ID = os.getenv("POSTGRES_CONN_ID", "postgres_default")
_POSTGRES_DATABASE = os.getenv("POSTGRES_DATABASE", "postgres")
_POSTGRES_SCHEMA = os.getenv("POSTGRES_SCHEMA", "public")
_POSTGRES_TRANSFORMED_TABLE = os.getenv(
    "POSTGRES_WEATHER_TABLE_TRANSFORMED", f"model_weather_data_{SEQUENCE_NAME}"
)

_SQL_DIR = f"{os.getenv('AIRFLOW_HOME')}/include/sql/asset_sequences/{SEQUENCE_NAME}"
```



```

# First asset definition

@asset(
    schedule="@daily", # see: https://www.astronomer.io/docs/learn/scheduling-in-airflow-for-options
    tags=[ "assets", "seq_1", "XCom" ], # add tags in the UI
    params={
        "coordinates": Param({ "latitude": 46.9481, "longitude": 7.4474}, type="object")
    },
)
def seq_1_start(context):
    yield Metadata(Asset("seq_1_start"), {
        "latitude": context["params"]["coordinates"]["latitude"],
        "longitude": context["params"]["coordinates"]["longitude"],
        "ts": context["ts"],
    },
)

```



```

@asset(
    schedule=[seq_1_start],
    tags=[ "assets", "seq_1", "XCom" ],
)
def seq_1_create_table(context):
    from jinja2 import Template

    hook = PostgresHook(postgres_conn_id=_POSTGRES_CONN_ID)

    with open(f" {_SQL_DIR}/create_table_if_not_exists.sql", "r") as f:
        sql_template = f.read()

        template = Template(sql_template)
        sql = template.render(
            params={ "schema": _POSTGRES_SCHEMA, "table": _POSTGRES_TRANSFORMED_TABLE }
        )

    hook.run(sql)

    # You can attach metadata to an asset event. This is useful for passing small
    # amounts of
    # data specific to a sequence.

    # See: https://www.astronomer.io/docs/learn/airflow-datasets/#attaching-information-to-an-asset-event

```

```
@asset(  
    schedule=[seq_1_start],  
    tags=["assets", "seq_1", "XCom"],  
    params={  
        "coordinates": Param({"latitude": 46.9481, "longitude": 7.4474}, type="object")  
    }, # Airflow params can add interactive options on manual runs. See: https://www.astronomer.io/docs/learn/airflow-params  
)  
  
def seq_1_extract(context):  
    """  
    Extract data from the Open-Meteo API  
  
    Returns:  
        dict: The full API response  
    """  
  
    import requests  
  
    url = os.getenv("WEATHER_API_URL")  
  
    # if the DAG is run manually use the time stamp and lat/long of the manual run  
    if str(context["dag_run"].run_type) == "DagRunType.MANUAL":  
        latitude = context["params"]["coordinates"]["latitude"]  
        longitude = context["params"]["coordinates"]["longitude"]  
        ts = context["ts"]  
  
        # if the DAG is run as part of the asset sequence use the lat/long and timestamp of  
        # the upstream asset event  
    else:  
        # You can retrieve metadata from the triggering asset event.  
        # See: https://www.astronomer.io/docs/learn/airflow-datasets/#retrieving-asset-information-in-a-downstream-task  
        metadata_upstream = context["triggering_asset_events"] [  
            Asset("seq_1_start")  
        ][0].extra  
  
        latitude = metadata_upstream["latitude"]  
        longitude = metadata_upstream["longitude"]  
        ts = metadata_upstream["ts"]  
  
    url = url.format(latitude=latitude, longitude=longitude)
```

```
response = requests.get(url)

# attach the timestamp of the DAG run to the asset event
yield Metadata(Asset("seq_1_extract"), {"ts": ts})

# returning a value from an @asset pushes it to XCom
return response.json()

@asset(
    schedule=[seq_1_extract],
    tags=["assets", "seq_1", "XCom"],
)
def seq_1_transform(context) -> dict:
    """
    Transform the data
    Args:
        api_response (dict): The full API response
    Returns:
        dict: The transformed data
    """
    # if the DAG is run manually use the time stamp of the manual run
    if str(context["dag_run"].run_type) == "DagRunType.MANUAL":
        dag_run_timestamp = context["ts"]
    # if the DAG is run as part of the asset sequence use the timestamp of the upstream
    # asset event
    else:
        # You can retrieve metadata from the triggering asset event.
        # See: https://www.astronomer.io/docs/learn/airflow-datasets/#retrieving-asset-information-in-a-downstream-task
        metadata_upstream = context["triggering_asset_events"][Asset("seq_1_extract")][
            0
        ].extra
        dag_run_timestamp = metadata_upstream["ts"]

    # To retrieve data pushed by an upstream asset perform a cross-dag xcom pull
    api_response = context["ti"].xcom_pull(
        dag_id="seq_1_extract",
```

```
task_ids=[“seq_1_extract”],  
key=“return_value”,  
include_prior_dates=True,  
)[0]  
  
time = api_response[“hourly”][“time”]  
  
transformed_data = {  
    “temperature_2m”: api_response[“hourly”][“temperature_2m”],  
    “relative_humidity_2m”: api_response[“hourly”][“relative_humidity_2m”],  
    “precipitation_probability”: api_response[“hourly”][  
        “precipitation_probability”  
    ],  
    “timestamp”: time,  
    “date”: [  
        datetime.strptime(x, “%Y-%m-%dT%H:%M”).date().strftime(“%Y-%m-%d”)  
        for x in time  
    ],  
    “day”: [datetime.strptime(x, “%Y-%m-%dT%H:%M”).day for x in time],  
    “month”: [datetime.strptime(x, “%Y-%m-%dT%H:%M”).month for x in time],  
    “year”: [datetime.strptime(x, “%Y-%m-%dT%H:%M”).year for x in time],  
    “last_updated”: [dag_run_timestamp for i in range(len(time))],  
    “latitude”: [api_response[“latitude”] for i in range(len(time))],  
    “longitude”: [api_response[“longitude”] for i in range(len(time))],  
}  
  
return transformed_data  
  
  


```
@asset(
 schedule=[seq_1_create_table, seq_1_transform],
 tags=[“assets”, “seq_1”, “XCom”],
)
def seq_1_load(context):
 """
 Load the data to the destination without using a temporary CSV file.
 Args:
 transformed_data (dict): The transformed data
 """
```


```

```
import csv
import io
from airflow.providers.postgres.hooks.postgres import PostgresHook

transformed_data = context["ti"].xcom_pull(
    dag_id="seq_1_transform",
    task_ids=["seq_1_transform"],
    key="return_value",
    include_prior_dates=True,
)[0]

hook = PostgresHook(postgres_conn_id=_POSTGRES_CONN_ID)

csv_buffer = io.StringIO()
writer = csv.writer(csv_buffer)
writer.writerow(WEATHER_COL_ORDER)
rows = zip(*[transformed_data[col] for col in WEATHER_COL_ORDER])
writer.writerows(rows)

csv_buffer.seek(0)

with open(f"({_SQL_DIR})/copy_insert.sql") as f:
    sql = f.read()
sql = sql.replace("{schema}", _POSTGRES_SCHEMA)
sql = sql.replace("{table}", _POSTGRES_TRANSFORMED_TABLE)

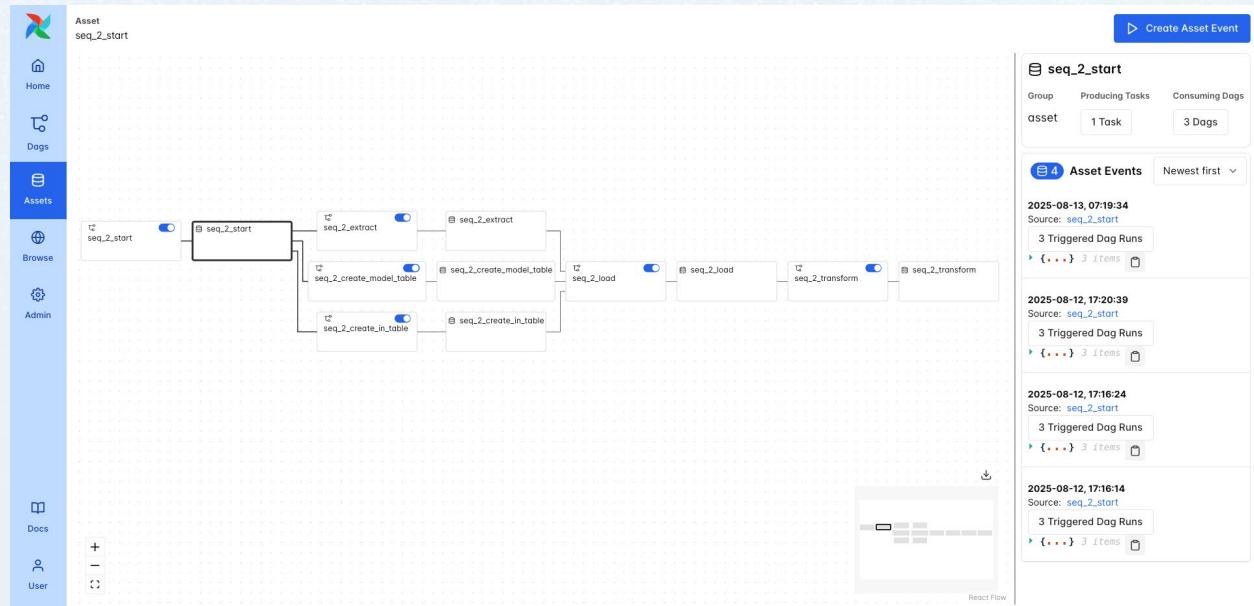
conn = hook.get_conn()
cursor = conn.cursor()
cursor.copy_expert(sql=sql, file=csv_buffer)
conn.commit()
cursor.close()
conn.close()
```

Asset sequence ETL using explicit external storage

This asset sequence accomplishes the same as the [ELT pattern DAG using explicit external storage](#), extracting data from the Open-Meteo API, loading the raw data into a Postgres database and then running a transformation on it, but without passing any data through XCom.

In the companion GitHub repository you can find this DAG [here](#) and its supporting SQL scripts [here](#).

This is the asset graph that is created:



```
"""
```

```
## Simple asset sequence creating 4 DAGs to perform an ELT pattern with intermediary storage
```

These 4 asset decorated functions created 4 DAGs each containing one task following a similar pattern as the regular `elt_intermediary_storage` DAG.

Note that passing the data between assets via XCom necessitates a Cross-DAG xcom pull and other information like DAG params and timestamps are passed via the metadata asset events.

```
"""
```

```
import os
import json

from airflow.providers.amazon.aws.hooks.s3 import S3Hook
from airflow.sdk import Param, asset, Metadata, Asset
from airflow.providers.postgres.hooks.postgres import PostgresHook

# ----- #
# DAG-level variables #
# ----- #

SEQUENCE_NAME = os.path.basename(__file__).replace(".py", "")

_AWS_CONN_ID = os.getenv("MINIO_CONN_ID", "minio_local")
_S3_BUCKET = os.getenv("S3_BUCKET", "open-meteo-etl")

_POSTGRES_CONN_ID = os.getenv("POSTGRES_CONN_ID", "postgres_default")
_POSTGRES_DATABASE = os.getenv("POSTGRES_DATABASE", "postgres")
_POSTGRES_SCHEMA = os.getenv("POSTGRES_SCHEMA", "public")
_POSTGRES_IN_TABLE = os.getenv("POSTGRES_WEATHER_TABLE_IN", f"in_weather_data_{SEQUENCE_NAME}")
_POSTGRES_TRANSFORMED_TABLE = os.getenv(
    "POSTGRES_WEATHER_TABLE_TRANSFORMED", f"model_weather_data_{SEQUENCE_NAME}"
```



```
sql_template = f.read()

template = Template(sql_template)
sql = template.render(
    params={"schema": _POSTGRES_SCHEMA, "table": _POSTGRES_IN_TABLE}
)

hook.run(sql)

@asset(
    schedule=[seq_2_start],
    tags=["assets", "seq_2", "XCom"],
)
def seq_2_create_model_table(context):
    from jinja2 import Template

    hook = PostgresHook(postgres_conn_id=_POSTGRES_CONN_ID)

    with open(f"{_SQL_DIR}/create_model_table_if_not_exists.sql", "r") as f:
        sql_template = f.read()

        template = Template(sql_template)
        sql = template.render(
            params={"schema": _POSTGRES_SCHEMA, "table": _POSTGRES_TRANSFORMED_TABLE}
        )

    hook.run(sql)

@asset(
    schedule=[seq_2_start],
    tags=["assets", "seq_2", "XCom"],
    params={
        "coordinates": Param({"latitude": 46.9481, "longitude": 7.4474}, type="object")
    },
)
```

```
def seq_2_extract(context):
    """
    Extract data from the Open-Meteo API
    Returns:
        dict: The full API response
    """
    import requests

    # if the DAG is run manually use the time stamp and lat/long of the manual run
    if str(context["dag_run"].run_type) == "DagRunType.MANUAL":
        latitude = context["params"]["coordinates"]["latitude"]
        longitude = context["params"]["coordinates"]["longitude"]
        dag_run_timestamp = context["ts"]

    # if the DAG is run as part of the asset sequence use the lat/long and timestamp of
    # the upstream asset event
    else:
        # You can retrieve metadata from the triggering asset event.
        # See: https://www.astronomer.io/docs/learn/airflow-datasets/#retrieving-asset-information-in-a-downstream-task
        metadata_upstream = context["triggering_asset_events"][Asset("seq_2_start")][
            0
        ].extra

        latitude = metadata_upstream["latitude"]
        longitude = metadata_upstream["longitude"]
        dag_run_timestamp = metadata_upstream["ts"]

    url = os.getenv("WEATHER_API_URL")

    url = url.format(latitude=latitude, longitude=longitude)

    response = requests.get(url).json()

    response_bytes = json.dumps(response).encode("utf-8")

    # Save the data to S3
    hook = S3Hook(aws_conn_id=_AWS_CONN_ID)
    hook.load_bytes(
```

```
bytes_data=response_bytes,  
key=f" {_INTERMEDIARY_STORAGE_KEY}/{dag_run_timestamp}.json",  
bucket_name=_S3_BUCKET,  
replace=True,  
)  
  
yield Metadata(Asset("seq_2_extract"), {"ts": dag_run_timestamp})  
  
  
@asset(  
    schedule=[seq_2_create_in_table, seq_2_create_model_table, seq_2_extract],  
    tags=["assets", "seq_2", "XCom"],  
)  
  
def seq_2_load(context):  
    """  
    Load the data from S3 to Postgres  
    """  
  
    # if the DAG is run manually use the time stamp of the manual run  
    if str(context["dag_run"].run_type) == "DagRunType.MANUAL":  
        dag_run_timestamp = context["ts"]  
  
    # if the DAG is run as part of the asset sequence use the timestamp of the upstream  
    # asset event  
    else:  
        # You can retrieve metadata from the triggering asset event.  
        # See: https://www.astronomer.io/docs/learn/airflow-datasets/#retrieving-asset-information-in-a-downstream-task  
        metadata_upstream = context["triggering_asset_events"][Asset("seq_2_extract")][  
            0  
        ].extra  
  
        dag_run_timestamp = metadata_upstream["ts"]  
  
    s3_hook = S3Hook(aws_conn_id=_AWS_CONN_ID)  
    response = s3_hook.read_key(  
        key=f" {_INTERMEDIARY_STORAGE_KEY}/{dag_run_timestamp}.json",  
        bucket_name=_S3_BUCKET,  
    )
```

```
api_response = json.loads(response)

postgres_hook = PostgresHook(postgres_conn_id=_POSTGRES_CONN_ID)

insert_sql = f"""
INSERT INTO {_POSTGRES_SCHEMA}.{_POSTGRES_IN_TABLE} (raw_data)
VALUES (%s::jsonb);
"""

postgres_hook.run(insert_sql, parameters=(json.dumps(api_response),))

@asset(
    schedule=[seq_2_load],
    tags=["assets", "seq_2", "XCom"],
)
def seq_2_transform(context):
    from jinja2 import Template

    hook = PostgresHook(postgres_conn_id=_POSTGRES_CONN_ID)

    with open(f" {_SQL_DIR}/transform.sql", "r") as f:
        sql_template = f.read()

    template = Template(sql_template)
    sql = template.render(
        params={
            "schema": _POSTGRES_SCHEMA,
            "in_table": _POSTGRES_IN_TABLE,
            "out_table": _POSTGRES_TRANSFORMED_TABLE,
        }
    )

    hook.run(sql)
```